

CJS in NIMBLE

NIMBLE 2022 virtual EFI workshop

NIMBLE Development Team

April 2022

Bonus example: CJS capture-recapture model with “classic” dipper data.

Dipper example:

- ▶ 294 Dippers monitored 1981-1987.
- ▶ One of the most classic capture-recapture teaching datasets ever.
- ▶ Thanks to Daniel Turek and Olivier Gimenez for Dipper examples from previous workshops.

Load the data

```
dipper_example_dir <- file.path("../", "../", "content", "example")
dipper <- read.csv(file.path(dipper_example_dir, "dipper.csv"))
y <- as.matrix(dipper[, 1:7])
y <- y + 1 # Code as 1 = not captured, 2 = captured.
first <- apply(y, 1, function(x) min(which(x != 1))) # first capture
y <- y[ first != 7, ] # remove records with first capture > 6
head(y)
```

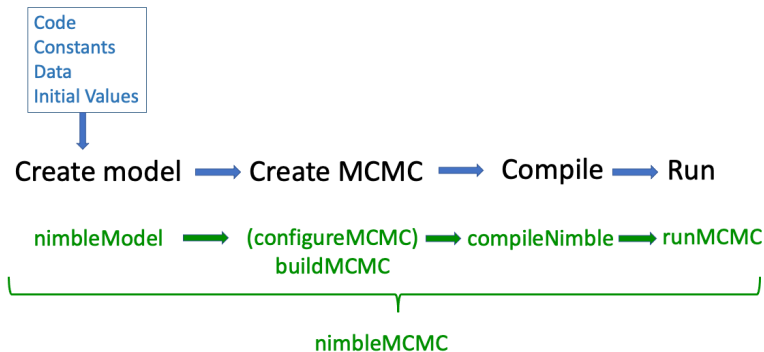
	year_1981	year_1982	year_1983	year_1984	year_1985	year_1986
## [1,]	2	2	2	2	2	2
## [2,]	2	2	2	2	2	2
## [3,]	2	2	2	2	1	1
## [4,]	2	2	2	2	1	1
## [5,]	2	2	1	2	2	2
## [6,]	2	2	1	1	1	1

Conventional Hidden Markov model code, slightly updated.

- ▶ data, y: 1 = not-detected, 2 = detected.
- ▶ latent states, z: 1 = alive, 2 = dead. (Following convention that “dead” is the last state.)
- ▶ Modified from Gimenez et al. capture-recapture workshop

```
dipper_code_dcat <- nimbleCode({
  phi ~ dunif(0, 1) # prior survival
  p ~ dunif(0, 1) # prior detection
  # likelihood
  gamma[1,1:2] <- c(phi, 1-phi) # Pr(alive t -> alive)
  gamma[2,1:2] <- c(0, 1) # Pr(dead t -> alive)
  delta[1:2] <- c(1, 0) # Pr(alive t = 1) = 1
  omega[1,1:2] <- c(1 - p, p) # Pr(alive t -> non-detected)
  omega[2,1:2] <- c(1, 0) # Pr(dead t -> non-detected)
  for (i in 1:N){
    z[i,first[i]] ~ dcat(delta[1:2]) # Illustrates initial
    for (j in (first[i]+1):T){
      z[i,j] ~ dcat(gamma[z[i,j-1], 1:2])
      y[i,j] ~ dcat(omega[z[i,j], 1:2])
    }
  }
})
```

Basic nimble workflow



Setup data, constants, and inits

```
zinit <- matrix(2, nrow = nrow(y), ncol = ncol(y)) # crea
zdata <- matrix(NA, nrow = nrow(y), ncol = ncol(y)) # crea
for(i in 1:nrow(zinit)) {
  known_alive <- range(which(y[i,] == 2))
  zinit[i, known_alive[1] : known_alive[2] ] <- NA # init
  zdata[i, known_alive[1] : known_alive[2] ] <- 1 # data
}
dipper_constants <- list(N = nrow(y),
                          T = ncol(y),
                          first = first)
dipper_data <- list(y = y,
                   z = zdata)
dipper_inits <- function() list(phi = runif(1,0,1),
                                p = runif(1,0,1),
                                z = zinit)
head(dipper_data$z) # data and inits have complementary
```

```
##      [,1] [,2] [,3] [,4] [,5] [,6] [,7]
```

```
## [1,] 1 1 1 1 1 1 NA
```

What are constants? What are data?

Constants are values needed to define model relationships

- ▶ Index starting/ending values like `N`
- ▶ Constant indexing vectors for indexing data groupings (site, treatment, individual, time): `beta[treatment[i]]`.
- ▶ Constants must be provided when creating a model with `nimbleModel`.

Data represents a flag on the role a node plays in the model

- ▶ E.g., data nodes shouldn't be sampled in MCMC.
- ▶ Data *values* can be changed.
- ▶ Data can be provided when calling `nimbleModel` or later.

Providing data and constants together.

- ▶ Data and constants can be provided together **as constants**.
- ▶ It would be slightly easier for BUGS/JAGS users to call this “data”, but that would blur the concepts.
- ▶ NIMBLE will usually disambiguate data and constants when they are provided together as constants.

What are covariates and other

Constants vs data for nested indexing

When values are grouped (particularly in irregular ways), we often have (potentially complicated) indexing.

Here is a model code snippet from a GLMM for disease occurrence in deer from different farms.

- ▶ sex is coded as 1 or 2 to index fixed-effect intercepts.
- ▶ farm_ids are coded as 1-24 for random-effect intercepts for data from 24 farms.
- ▶ Membership indices that are known in advance.
- ▶ **Provide known indexing vectors in constants.**
- ▶ Otherwise nimble allows that you (or your model) might change them later, which makes it handle them inefficiently (with unnecessary computation).

```
nimbleCode({  
  # ...incomplete code snippet...  
  for(i in 1:2) sex_int[i] ~ dnorm(0, sd = 1000) # prior  
  for(i in 1:num_farms) farm_effect[i] ~ dnorm(0, sd = farm  
  farm sd ~ dunif(0, 100) # prior
```

Create a model

```
dipper_model <- nimbleModel(code = dipper_code_dcat,  
                             constants = dipper_constants,  
                             data = dipper_data,      # data  
                             inits = dipper_inits()   # inits  
                             )                        # dimensions
```

```
## Defining model
```

```
## Building model
```

```
## Setting data and initial values
```

```
## Running calculate on model
```

```
## [Note] Any error reports that follow may simply reflect
```

```
## Checking model sizes and dimensions
```

Create an MCMC

```
dipper_MCMCconf <- configureMCMC(dipper_model, monitors = c(

## ===== Monitors =====
## thin = 1: p, phi
## ===== Samplers =====
## RW sampler (2)
##   - phi
##   - p
## categorical sampler (613)
##   - z[] (613 elements)

dipper_MCMC <- buildMCMC(dipper_MCMCconf)
```

Compile the model and MCMC

3. Compile the model and MCMC

```
C_dipper_model <- compileNimble(dipper_model) # These two  
  
## Compiling  
## [Note] This may take a minute.  
## [Note] Use 'showCompilerOutput = TRUE' to see C++ comp  
  
C_dipper_MCMC <- compileNimble(dipper_MCMC, project = dipper  
  
## Compiling  
## [Note] This may take a minute.  
## [Note] Use 'showCompilerOutput = TRUE' to see C++ comp
```

4. Run the MCMC

```
samples <- runMCMC(C_dipper_MCMC, niter = 10000, samplesAsC  
  
## Running chain 1 ...  
## |-----|-----|-----|-----|  
## |-----|-----|-----|-----|  
  
# Alternative:
```

How can I use the model in R?

```
## Defining model

## Building model

## Setting data and initial values

## Running calculate on model
## [Note] Any error reports that follow may simply reflect

## Checking model sizes and dimensions

## Compiling
## [Note] This may take a minute.
## [Note] Use 'showCompilerOutput = TRUE' to see C++ compiler output

class(dipper_model)[1] # This is a reference class (S5) object

## [1] "dipper_cod_MID_2_modelClass_UID_17_UID_18"

dipper_model$gamma # Look at a model variable,

##           [,1]           [,2]
```

NIMBLE might insert nodes into your model!

These are called *lifted nodes*.

Example 1: reparameterization

You give NIMBLE this:

```
nimbleCode({  
  tau <- 1E-0.6  
  mu ~ dnorm(0, tau)  
})
```

- ▶ NIMBLE defaults to parameterizations from **WinBUGS/OpenBUGS/JAGS, not R.**
- ▶ Default SD/Var/precision for dnorm is **precision** = 1/variance.
- ▶ NIMBLE converts this to a *canonical* parameterization for computations by treating it like this:

```
nimbleCode({  
  tau <- 1E-0.6  
  some_long_name_created_by_nimble <- 1/sqrt(tau) # a lifted node  
  mu ~ dnorm(0, sd = some_long_name_created_by_nimble)  
})
```

How can I use the MCMC configuration in R?

- ▶ Change the set of samples that compose an MCMC algorithm.
- ▶ See `help(samplers)` for samplers built in to `nimble`.
- ▶ Default sampler assignments:
 - ▶ Conjugate (Gibbs) sampler when possible.
 - ▶ Special samplers for Bernoulli, categorical, Dirichlet, multinomial, possibly others.
 - ▶ Slice samplers for discrete distributions such as Poisson.
 - ▶ Adaptive random-walk Metropolis-Hastings samplers for other continuous distributions.
- ▶ $\text{MCMC efficiency} = \frac{\text{Effective sample size (mixing)}}{\text{computation time}}$
 - ▶ Both speed and mixing matter.
 - ▶ There is often a tradeoff between these.
- ▶ Some ways to customize samplers
 - ▶ Block (jointly) sample correlated dimensions.
 - ▶ Block (joint) samplers include adaptive random-walk

How can I use uncompiled vs. compiled models and algorithms?

- ▶ An important and perhaps unfamiliar principle:
 - ▶ (Almost) everything can be run **uncompiled** (in R) or **compiled** (in C++).
- ▶ Uncompiled use of models and algorithms in R allows debugging.
 - ▶ Behavior is not always identical but is close.
 - ▶ Example: Error trapping will behave differently. Errors in C++ might not occur in R.

Differences from NIMBLE to JAGS and/or BUGS

- ▶ See our guide on converting from JAGS to NIMBLE
- ▶ Wrap your model code in `nimbleCode({})` directly in R.
- ▶ Provide information about missing or empty indices:
 - ▶ Use `x[1:n, 1:m]` or
 - ▶ `x[,]` with `dimensions = list(x = c(n,m))`.
 - ▶ Do not use `x` without brackets (unless it is a scalar).
- ▶ Decide how much control you need:
 - ▶ `nimbleMCMC` will do everything.
 - ▶ `nimbleModel`, `configureMCMC`, `buildMCMC`, `compileNimble` and `runMCMC` give you more control.